



**Question 1. Concurrency** [4 MARKS]

Suppose two threads execute the following C code concurrently, accessing shared variables *a*, *b*, and *c*:

```
// initialization code
int a = 4;
int b = 0;
int c = 0;

// Thread 1                                // Thread 2

if (a < 0) {                                b = 10;
    c = b - a;                               a = -3;
} else {
    c = b + a;
}
```

What are the possible values for *c* after both threads complete? You can assume that reads and writes of the variables are atomic, and that the order of statements within each thread is preserved in the code generated by the C compiler. Marks will be deducted for wrong or too many answers.

13

4

14

7

**Question 2. Synchronization** [8 MARKS]

To make laughing gas, two Nitrogen (N) atoms and one Oxygen (O) atom have to be fused. Let us synthesize this reaction by treating the atoms as threads and synchronizing them.

Each N atom invokes a procedure `nReady()` when it is ready to react, and each O atom invokes a procedure `oReady()` when it is ready. The procedures delay until there are at least two N atoms and one O atom present, and then the `oReady()` procedure calls a `Laugh()` procedure. After the `Laugh()` call, the two instances of `nReady()` and one instance of `oReady()` return.

To be fair, the N and the O atoms must fuse in the order they arrive. For example, N1 and N2 should fuse with O1, N3 and N4 should fuse with O2, and so on.

**N** → The code below shows a proposed solution for creating laughing gas. The `count` variable tracks the order in which the O atoms arrive. You may assume that the semaphore implementation enforces FIFO order for wakeups, i.e., the thread waiting longest in `P()` is always the next thread woken up by a call to `V()`.

```
Semaphore nwait = 0;
Semaphore owait = 0;
int count = 0;

nReady()
{
    count++;
    if (count % 2 == 1) {
        P(nwait);
    } else {
        V(owait);
        P(nwait);
    }
    return;
}

oReady()
{
    P(owait);
    Laugh();
    V(nwait);
    V(nwait);
    return;
}
```

**Part (a)** [2 MARKS] Unfortunately, this code has a race. Describe a case in which this code will not work.

The race is in the update of `count`. Multiple N atoms share `count`, and so `count` needs to be updated in a critical section.

**Part (b)** [6 MARKS] Below, you will try to fix the race. We have added the `lock()` code below. You need to add `unlock()` in the code below. Write within each box shown below, **one** of the following: 1) **unlock**, if the code needs an unlock there, 2) **deadlock**, if adding an unlock there may cause a deadlock, and 3) **incorrect**, if adding an unlock there may cause any other kind of incorrect synchronization.

```
Semaphore hwait = 0;
Semaphore owait = 0;
int count = 0;
```

```
Lock L;    // new code
```

```
nReady()
{
    lock(L); // new code
    count++;
```

```
    if (count % 2 == 1) {
```

```
        P(nwait);
```

```
    } else {
```

```
        V(owait);
```

```
        P(nwait);
```

```
    }
```

```
    return;
```

```
}
```

⇐ deadlock: 3 N threads update count to 3, and then all will get stuck in P(nwait)

⇐

⇐ deadlock: after 1st N thread waits at P(), no one else gets to run.

⇐ incorrect: N thread 2 and 3 could fuse with O thread 1. However, we also accepted an unlock here as acceptable answer.

⇐

⇐ deadlock: after 1st N thread waits at P(), no one else gets to run.

**Question 3. Scheduling** [10 MARKS]

There are three processes, A, B and C in the system. Process A is CPU-bound and always runnable. Processes B and C are IO-bound. After Process B runs for 1 unit of time, it needs to block for 1 unit of time. After Process C runs for 1 unit of time, it needs to block for 3 units of time.

**Part (a)** [6 MARKS] Show how the three processes are scheduled by a round-robin scheduler. The time slice of the scheduler is 1 unit of time. The first 12 time slices are shown in the three tables below. Show the process that is running in the “Running Process” table. Also, show the runnable processes in the “Ready Queue” table, and the blocked processes in the “Wait Queue” table, at each time slice.

Start by scheduling Processes A, B, and C, as shown below. Assume that a process becomes runnable just before the arrival of a timer interrupt. Also, assume that if multiple processes become runnable at the same time, then they are added to the ready queue in the order A, B, and C.

Running Process											
1	2	3	4	5	6	7	8	9	10	11	12
A	B	C	A	B	A	B	C	A	B	A	B

Ready Queue											
head	B	C	A	B	A		C	A	B	A	C
	C	A					A				A
tail											

Wait Queue											
head		B	C	C	C		B	C	C	C	
					B					B	
tail											

If there is a repeating pattern of running processes, circle the pattern in the “Running Process” table.

Roughly, what fraction of the CPU does each process take?

A = Based on repeating pattern, over the long term,  $2/5$

B =  $2/5$

C =  $1/5$

We also accepted  $5/12$ ,  $5/12$ , and  $2/12$  as answers based on the first 12 time units.

**Part (b)** [4 MARKS] Show how the three processes are scheduled by a fixed priority scheduler. Assume Process C runs with the highest priority, Process B with medium priority, and Process A with the lowest priority. Start by scheduling Process C, as shown below.

Running Process												
1	2	3	4	5	6	7	8	9	10	11	12	
C	B	A	B	C	B	A	B	C	B	A	B	

Ready Queue												
head	B	A		A	A	A		A	A	A		A
	A											
tail												

Wait Queue												
head		C	C	C	B	C	C	C	B	C	C	C
			B				B				B	
tail												

If there is a repeating pattern of running processes, circle the pattern in the “Running Process” table.

Roughly, what fraction of the CPU does each process take?

A = Based on repeating pattern, over the long term,  $1/4$

B =  $1/2$

C =  $1/4$

These same numbers also apply for the 12 time units.

**Question 4. Paging [6 MARKS]**

The figure below shows the address translation mechanism used by a 64-bit processor. Starting with a 48-bit virtual addresses, it uses 4 levels of page table to translate the address to a 52-bit physical address. Each page table entry is 8 bytes long.




Figure 1: 4-level Page Table

If the page size is increased, the number of levels of page table can be reduced. How large must pages be in order to translate 48-bit virtual addresses with only a 2 level page table? Draw the address translation mechanism, similar to the figure above, corresponding to this page size.

Notice first that the page size is  $2^{12}$  in Figure 1. The page table size is also  $2^{12}$  ( $2^9 \times 8$  bytes per page table entry, for a 64 bit processor).

For a two level page table, say page size is  $2^n$ . The page table size is also  $2^n$  bytes in this multi-level page table. So the number of PTE in a page table is  $2^{(n-3)}$ .

So the virtual address is broken into  $(n-3)$ ,  $(n-3)$ , and  $n$  bits.

$$(n-3) + (n-3) + n = 48$$

$n = 18$ , so page size is  $2^{18} = 256\text{KB}$ .

The 48-bit virtual address is divided into 15, 15, 18 bits. The first 15 bits index into the first page table, the second 15 bits index in the second page table, and the rest of the 18 bits is the page offset.

**Question 5. TLB [8 MARKS]**

A 32-bit processor uses a two-level page table, with a page size of 4KB. It has two TLBs, an I-TLB (instruction TLB) that stores mappings for code addresses, and a D-TLB (data TLB) that stores mappings for data addresses. Both TLBs are fully associative (all TLB entries are looked up in parallel), and both have 64 entries. Consider the following code snippet:

```
char array[4096 * 64];

void simple() {
    int i = 0;
    int or = 0;
    int and = 1;

    for (i = 0; i < 4096 * 64; i++) {
        or = or | array[i];
    }
    for (i = 0; i < 4096 * 64; i++) {
        and = and & array[i];
    }
}
```

**Part (a)** [1 MARK] How many I-TLB misses will take place when the `simple()` function is run? Provide a justification for each TLB miss. State any assumptions that you make.

One I-TLB miss, assuming all the code above fits in 4KB.

**Part (b)** [3 MARKS] What is the **minimum** number of D-TLB misses that will take place when the `simple()` function is run? Provide a justification for each TLB miss. State any assumptions that you make.

Minimum number of D-TLB misses = 1 + 63 + 1 + 1 (see reasons below) = 66

Notice that array is stored in 64 pages

- 1 D-TLB miss for a stack page (which stores the local variables of `simple()`).

Loop 1:

- 63 D-TLB misses when the first 63 pages of the array are accessed. Now the D-TLB is full.  
- 1 D-TLB miss when the last element of array is accessed. This access will require invalidating a TLB entry. Let's invalidate the last TLB entry (for the 63th page), and replace it with the TLB entry for the 64th page.

Loop 2:

- 0 D-TLB misses for the first 62 page accesses (these entries are already present).  
- 1 D-TLB miss for page 63. Let's invalidate TLB entry for page 62, and replace it with TLB entry for page 63.  
- 0 D-TLB miss for page 64.



**Part (c)** [2 MARKS] Let us estimate the performance overhead of using the paging system. Assume that we have disabled the processor cache, so that every load and store instruction accesses memory (RAM). Assume also that there is no cost to accessing the TLB, and the compiler has performed no optimizations. Also, we will only consider accesses to `array` in the `simple()` function. How many memory accesses are performed to read `array`? How many memory accesses are performed to read the page table? Now estimate the overhead of the paging system. State your answer as a single number, possibly in terms of a power of 2.

Total number of memory access for reading array:  $4096 * 64 * 2$

Number of memory accesses to read the page table =  $2$  (nr. of levels of page table) \* nr. of TLB misses =  $2 * 66$  (from previous answer)

$$\begin{aligned} \text{Overhead} &= 2 * 66 / (2 * 66 + 2 * 64 * 4096) \\ &\sim (2 * 64) / (2 * 64 * 4096) \\ &\sim 1/4096 \\ &= 2^{(-12)} \end{aligned}$$

If you said  $2^{(-11)}$ , we accepted the answer. This is reasonable if you simply assume that there is one TLB miss for each page. So you require two additional memory accesses for the page table for accessing each page (4096 memory accesses).

**Part (d)** [2 MARKS] Suggest a way of changing the code in the `simple()` function to reduce the number of D-TLB misses. What is the number of misses with your code change? Provide a justification.

Use one loop to do both calculations.

This reduces the one additional D-TLB miss that occurred in Part(b) when loop 2 executed.

Total number of D-TLB misses = 65

*[Use the space below for rough work.]*