

Question 1. All False [5 MARKS]

Your friend makes the following statements below. You know they are not correct. Clearly explain to your friend **why** they are incorrect. We have provided an example below. Single sentence answers are preferred.

0) On a uni-processor, when a process makes a system call, the system call must return to the process before another process can make a system call.

Answer: A system call can block in the kernel, after which another process can run and issue a system call before the first process's system call returns.

A) When a process invokes a system call, a context switch occurs.

B) A divide-by-zero error in a program invokes the trap instruction.

C) A program never executes code after the Unix `execve()` system call.

D) A Unix `wait()` system call will always cause a process to block.

E) A `setcontext()` call will always enable signals.

Question 2. Thread Switching [5 MARKS]

You know that implementing thread switching is a little tricky. To simplify matters, you start by implementing thread switching between just two threads, A and B, as shown below. You have added A and B to the run queue, and initialized their thread context so that they will start executing the `thread_A()` and `thread_B()` functions when they are run for the first time. Your scheduler runs Thread A first.

```
int i = 0;
ucontext_t uA, uB;

thread_A() {
    int d = 0;
    while (i < 3) {
        i++;
        printf("A:%d_", i);
        d = 0;
        getcontext(&uA);
        if (d == 0) {
            d = 1;
            setcontext(&uB);
        }
    }
}

thread_B() {
    int d = 1;
    while (i < 3) {
        i++;
        printf("B:%d_", i);
        d = 1;
        getcontext(&uB);
        if (d == 1) {
            d = 0;
            setcontext(&uA);
        }
    }
}
```

Part (a) [3 MARKS] Circle the output you expect to see? Hint: Think carefully about what is saved by the `getcontext()` function.

- A) A:1 A:2 A:3
- B) A:1 B:2 B:3
- C) A:1 B:2 A:3 B:4 A:5 B:6
- D) A:1 B:2
- E) A:1 B:1
- F) A:1 B:2 A:3

Part (b) [2 MARKS] Briefly explain why you have chosen the answer above.

Question 3. Synchronization [5 MARKS]

Consider the following program that uses three threads (A, B and C). It synchronizes these threads using locks and condition variables. The program begins by invoking the `thread_A()` function.

```
struct lock *l;
struct cv *cv;

void thread_A()
{
    l = lock_create();
    cv = cv_create();
    thread_create(thread_B, 0);
    thread_create(thread_C, 0);
    thread_yield(THREAD_ANY);
    thread_yield(THREAD_ANY);
    printf("STOP_HERE\n");
}

void thread_B(void *arg)
{
    lock_acquire(l);
    cv_signal(cv, l);
    thread_yield(THREAD_ANY);
    cv_wait(cv, l);
    lock_release(l);
}

void thread_C(void *arg)
{
    lock_acquire(l);
    cv_wait(cv, l);
    thread_yield(THREAD_ANY);
    cv_signal(cv, l);
    lock_release(l);
}
```

Part (a) [3 MARKS] Trace the execution of this program until it prints out the message “STOP HERE” by writing down the sequence of context switches that have occurred up to this point. Assume that the scheduler runs threads in FIFO order with no time-slicing (non-preemptive scheduling), all threads have the same priority, and threads are placed in wait queues in FIFO order. The output shown below should be in the form $B \rightarrow A \rightarrow C$, signifying that thread B context switches to thread A, which then context switches to thread C. Hint: Think carefully about how the condition variable primitives work.

Part (b) [2 MARKS] When the program prints out the message “STOP HERE”, list the currently running thread, and the (zero or more) threads in the ready and the wait queues shown below.

current thread:

ready queue:

lock wait queue:

cv wait queue:

Question 4. Waiting for Exit [5 MARKS]

You have implemented an early version of the Unix operating system that provides the signal facility (e.g., the `kill` system call for sending signals). Your operating system also provides the `sleep` system call, which blocks a process until a signal is sent to the process.

Your operating system, however, does not provide the `wait` system call to allow a parent to wait until its child process exits. You consider implementing the functionality of the `wait` system call at the user level with the code shown below. The initial call to `signal` registers an empty signal handler. The `sleep` call blocks until a signal is delivered, after which it returns.

```
void
null(int unused)
{
}

int
main()
{
    int pid;

    signal(SIGUSR1, null);
    pid = fork();
    if (pid) {
        /* block until child sends signal */
        sleep();
        printf("child_is_dead\n");
        exit(0);
    } else {
        kill(getppid(), SIGUSR1);
        exit(0);
    }
}
```

Describe two reasons why the code above does not implement the Unix `wait` functionality correctly.

1)

2)

Question 5. Scheduling [5 MARKS]

Consider the five threads shown in the table below. Their arrival times and their processing times are known and shown below.

Thread	Arrival Time	Processing Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

A thread scheduler runs these threads in the order shown below. The scheduler has a time slice of 1 time unit. You can assume that threads arrive just before the time slice expires.

Time	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19
Thread	1	1	2	2	3	2	4	4	3	2	4	5	3	2	4	1	5	3	2	4

Part (a) [1 MARK] Circle either a) or b)

- a) Scheduler is non-preemptive
- b) Scheduler is preemptive

Part (b) [3 MARKS] Briefly describe the policy that you think is being implemented by the scheduler.

Part (c) [1 MARK] Describe one problem that can occur with this type of scheduling.

[Use the space below for rough work.]