

Duration: 1 hour 15 min

Examiner: A. Goel

Please fill your student number and name below and then read the instructions below carefully.

Student Number: _____

First Name: _____

Last Name: _____

Instructions

Examination Aids: No examination aids are allowed.

Do not turn this page until you have received the signal to start.

Do not remove any sheets from this test book. Answer all questions in the space provided. No additional sheets are permitted. Use the last blank page as scratch space. This page will not be marked.

This exam consists of 5 questions on 8 pages (including this page). The value of each part of each question is indicated. The total value of all questions is 30.

For the written answers, explain your reasoning clearly. Be as brief and specific as possible. Clear, concise answers will be given higher marks than vague, wordy answers. Marks will be deducted for incorrect statements in an answer. Please write legibly!

Work independently.

MARKING GUIDE

Q1: _____ (5)

Q2: _____ (8)

Q3: _____ (9)

Q4: _____ (6)

Q5: _____ (2)

TOTAL: _____ (30)

Question 1. Labs Are Fun [5 MARKS]

No one wants to hold hot potatoes. The code shown below is invoked by each thread. It tries to move potatoes between participant threads, arranged in a circle. The code is designed to ensure that only *one* thread has a potato at a time. There is a single `potato_lock` global variable.

```
1 int potato[NTHREADS] = {1, 0}; /* potato[0] = 1, potato[i] = 0 when i > 0 */
2
3 void try_move_potato(int num) /* num is the thread number */
4 {
5     assert(interrupts_enabled());
6     lock(potato_lock);
7     if (potato[num]) {
8         potato[num] = 0;
9         potato[(num + 1) % NTHREADS] = 1; /* pass potato to next thread */
10    }
11    unlock(potato_lock);
12 }
```

Part (a) [2 MARKS] Do we need the `lock` and `unlock` code shown above to ensure that only one thread has a potato at a time? If so, show an interleaving that can cause a problem. If not, explain why.

Part (b) [3 MARKS] Suppose we switch lines 8 and 9 (shown in bold) in the code above. Now do we need the `lock` code to ensure that only one thread has a potato at a time? Again, explain your answer.

Question 2. System Calls [8 MARKS]

When a process invokes the `fork` system call, a child process is created. This process is a copy of the calling process. As we have seen in class, the parent and the child processes run in a unsynchronized manner. In this problem, we will aim to synchronize the two processes. Consider the program listing below. The parent needs to wait until the child has been initialized.

```
int
main()
{
    int pid;
    int child_is_initialized = 0;

    pid = fork();
    if (pid) { /* parent process */
        while (child_is_initialized == 0) { /* spin */
        }
        ...
    } else { /* child process */
        /* initialize child */
        /* let the parent know that child has been initialized */
        child_is_initialized = 1;
        ...
    }
}
```

Part (a) [1 MARK] Briefly explain why the code above does not work correctly.

Part (b) [1 MARK] Briefly explain why this synchronization cannot be performed without using system calls.

Part (c) [4 MARKS] You decide to implement two system calls, `down()` and `up()`, similar to semaphores, to perform the synchronization above. The parent invokes the `down()` call, and the child invokes the `up()` call.

Show what arguments you would pass to these system calls. Make sure to explain why you are passing these arguments. If you need to use new variables in your code, make sure to declare them below (including their types). Assume that you are only allowed to use these two calls, and they implement synchronization correctly in the kernel.

```
down (
```

```
) ;
```

```
up (
```

```
) ;
```

Part (d) [2 MARKS] Can you think of a way that the kernel can ensure that another process invoking `down()` or `up()` will not interfere with the synchronization above.

Question 3. Synchronization [9 MARKS]

Consider a calculator program that supports arithmetic operations and variable assignments (e.g., $a = 5$, $b = a + 7$). The variable names can only be a single character.

You want to support concurrent computations using condition variables. You want two computations to run concurrently if they don't use any common variables. However, each computation must run atomically. For example, if a computation is using variables a and b , then other computations using *either* a *or* b should wait. You make a function `get_vars(const char *expr, char *vars)` that returns the number and the list of all the variables that are used in the `expr` expression. For example, `get_vars("`b = a + 7'", vars)` returns 2, since 2 variables are used. It also sets the first 2 characters of the `vars` array to a and b .

In the following function, fill in the blank spaces to make your concurrent calculator. All spaces may not need to be filled. You may only use one lock and one condition variable, as declared. You may use `wait(cv, lock)`, `signal(cv, lock)` and `broadcast(cv, lock)`.

```
int symbol_state[256] = { 0 }; /* 256 ascii chars, all elements are 0 */
struct mutex * lock;          /* assume lock and cv are initialized */
struct cv     * cv;
int calculate(const char *expr, int len) { /* assume expr is correct */
    char vars[256];
    int num_vars = get_vars(expr, vars); /* return chars in expr in vars */
    int done = 0, answer, i;

    while (!done) {
        done = 1;
        for (i = 0; i < num_vars; ++i)
            if (symbol_state[vars[i]] == 1) {
                break; /* breaks out of for loop, not while loop */
            }
    }

    /* make sure to not run computation within a lock */
    answer = compute_expression(expr, len);

    for (i = 0; i < num_vars; ++i)
        symbol_state[vars[i]] = 0;

    return answer;
}
```

Question 4. Scheduling [6 MARKS]

Consider the 8 threads shown below, with their associated arrival times and processing times:

Job	Arrival Time	Processing Time	Job	Arrival Time	Processing Time
A	0	5	E	7	3
B	2	2	F	8	4
C	3	6	G	9	1
D	5	1	H	11	2

For the two scheduling policies shown below, show the sequence of threads that run on a single processor (if a thread runs for multiple consecutive time units, show it once). When there is a tie, assume FIFO scheduling. For preemptive policies, assume a time slice of *2 time units*. If a job finishes without using up its time slice, scheduling occurs immediately. Also assume that threads arrive just before the time slice expires.

As an example, with FIFO, your output would be:

A	B	C	D	E	F	G	H
---	---	---	---	---	---	---	---

Part (a) [3 MARKS] Longest Job First

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

Part (b) [3 MARKS] Shortest Remaining Time

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

Question 5. Labs Really Are Fun [2 MARKS]

You know that implementing thread switching is a little tricky. To simplify matters, you start by implementing thread switching between just two threads, A and B, as shown below. Both have the same code, other than the arguments to the calls to `getcontext`, `setcontext` and the `printf` functions. Assume that the code is run on a single processor.

```
ucontext_t uA, uB;

thread_A() {
    int done = 0;

    while (1) {
        getcontext(&uA);
        if (done == 0) {
            done = 1;
            setcontext(&uB);
        }
        done = 0;
        printf("A_makes_progress\n");
    }
}

thread_B() {
    int done = 0;

    while (1) {
        getcontext(&uB);
        if (done == 0) {
            done = 1;
            setcontext(&uA);
        }
        done = 0;
        printf("B_makes_progress\n");
    }
}
```

You find that when you compile this code normally, it shows the output that you expect: `A makes progress` and `B makes progress` are printed alternately. Now you want to run this code fast, and so you compile this code with compiler optimization options turned on (e.g., `gcc -O thread.c`). When you run the optimized code, you see no output!

When you add some more debugging output in your code, you realize that the context switching is happening. Based on what you have learned about context switching, what might the compiler be doing so that the `printf()` in both the threads shown above does not execute? Assume that the compiler optimizations are generating correct code.

[Use the space below for rough work.]